

Building a Telugu + Indian-English TTS Dataset

Sarvam AI take-home

The short version

I set out to build a 60-minute, single-speaker text-to-speech dataset – roughly half Telugu, half Indian English – where every clip is one person speaking cleanly, with an accurate transcript and a plain-English sentence describing *how* it's said (the speaker's emotion, style, pace, pitch). I collected a lot more than I needed from 14 YouTube sources, ran it through a quality pipeline, listened to the data (my grandmother helped with the Telugu), and trimmed it down to a balanced set.

Where it landed: 190 clips, 79 minutes, split evenly (95 Telugu / 95 English). This is a deliberately over-sized "review buffer" – bigger than 60 minutes on purpose, so clips can be thrown out during review and the set still clears the 30-min-per-language bar. The script `pipeline/balance.py` re-trims it to the final size whenever I want.

The assignment says it's graded on *data quality and judgment, not code*, so most of this report is the story: what I built, the things that went wrong, and what I'd fix.

A quick glossary

A few terms show up throughout. In plain language:

Term	What it means here
ASR	Automatic Speech Recognition – turning audio into text. Sarvam's model is <code>saaras:v3</code> .
VAD	Voice Activity Detection – finding where speech is vs. silence, so I can cut clips at pauses.
Diarization	"Who spoke when" – splitting a recording by speaker, so I can keep just one person.
SER	Speech Emotion Recognition – a model that guesses emotion from the <i>sound</i> , not the words.
Parler-style description	A natural-language sentence describing delivery (e.g. "A woman narrates slowly in a calm voice"), the format the Parler-TTS project popularised.

Term	What it means here
Code-switching	Mixing two languages in one sentence – extremely common in Indian speech (Telugu with English words).
WER / CER	Word / Character Error Rate – how far a transcript is from the correct one (lower is better).
Manifest	My single bookkeeping file, data/manifest.jsonl – one row per clip holding everything known about it.

The sources I collected (what the names mean)

The dataset is only as good as what goes in, so I picked sources for *variety* – different speakers, genders, and speaking styles – not just whatever was easy. Each gets a short code-name in the config:

Source name	What it actually is	Lang	Why I chose it
Harshaneeyam	A Telugu literary storytelling podcast	te	emotional range (sad, excited)
BV_Pattabhiram	A Telugu motivational / counselling lecture	te	calm, serious delivery
Garikapati	Garikapati N. Rao's Telugu literary discourse	te	oratorical, animated
Te_Standup_M	A Telugu stand-up comedy set (male)	te	playful/happy + a crowd-noise test
Te_ASMR	Telugu ASMR (whispered)	te	the only way to get <i>whisper</i> clips
En_Podcast_Raj	A long podcast that turned out to be Telugu (see iteration 6)	te	conversational; added a female Telugu voice
TED_India	A TEDx India talk	en	clear oratorical English
EN_Audiobook	An Indian-English audiobook narration (female)	en	clean narrative reference
EN_Podcast_Solo	A solo Indian-English podcast monologue (male)	en	conversational

Source name	What it actually is	Lang	Why I chose it
En_Comedy	"A Moment of Silence", a multi-host comedy podcast (female-led)	en	playful/happy; diarized to one speaker
En_Speech	An Indian-English speech	en	oratorical
En_Standup_F1/F2/F3	Three female stand-up specials	en	fix a male-heavy roster; happy/excited

So when I say "En_Comedy", that's a real comedy podcast with several hosts talking over each other – which is exactly why it needed diarization to pull out a single voice.

1. What I built and how the pipeline works

The whole thing is a chain of small scripts, each doing one job and writing its result back to one file (`data/manifest.jsonl`). Because every clip's progress is recorded in that file, I can stop and restart anywhere – which mattered, because I hit crashes and ran out of API credits more than once, and every restart just picked up where it left off.

Manifest (data/manifest.jsonl) is the single source of truth; every stage reads and writes it.




Figure 1: The pipeline. Each stage reads the manifest, does its one job, and writes back. Teal stages call a Sarvam API; the rest are local helpers. Rejected clips are kept (as an audit log), and a human-review step produces the trusted "gold" clips before export.

Walking through the stages:

- **Download** (s1) – grab each source's audio with `yt-dlp`, convert to 16 kHz mono and even out the loudness with `ffmpeg`.
- **Segment** (s2) – cut each recording into ~25-second clips, but *only at silences*, so a word is never chopped in half. Every clip stays under the 30-second limit Sarvam's ASR needs. For sources with more than one speaker, s2b instead runs Sarvam's diarization and keeps only the person who talks the most.
- **Filter** (s3) – listen for problems and *throw clips out*: background music, crowd noise, or two people in one clip. I never try to "repair" audio – for TTS, a cleaned-up clip is worse than no clip.
- **Features** (s4, s4b) – measure pitch, energy, and pace; and separately, run a model that guesses emotion straight from the *sound*.

- **Transcribe** (s5) – send each clip to Sarvam saaras:v3, telling it the language up front for accuracy.
- **Tag & describe** (s6, s7) – ask Sarvam’s sarvam-30b for the emotion, style, and whisper flag, then for the one-sentence description.
- **Export** (s8) – package it as a HuggingFace dataset and publish.

The one rule I’m proud of

The machine’s guesses and the human’s corrections are kept in **separate columns that never overwrite each other**. The transcript Sarvam produced lives in `asr_transcript`; if a human fixes it, that goes in `human_transcript`. They’re combined only at the very end:

```
final_transcript = human_transcript or asr_transcript    # human wins if present
final_emotion    = human_emotion    or llm_emotion
```

This means the raw machine output is never lost, so I can always measure honestly how good the ASR was. Human-checked clips become the trusted **gold** set. A manifest row, grouped by who owns each field:

```
identity : clip_id, source_channel, language, gender, duration, ...
machine  : asr_transcript / llm_emotion / llm_style / whisper / audio_emotion ...
human    : human_transcript / human_emotion / human_verified / reviewer ...
```

2. Iterations I did to improve data quality

Almost every quality gain came from catching the pipeline quietly doing the wrong thing. In order:

1. **The music filter wasn’t filtering anything.** The library I used (`inaSpeechSegmenter`) crashed on every single clip due to a version mismatch – but the crash was swallowed, so the stage “succeeded” while detecting zero music. I had to pin an older Keras and patch a NumPy call. *Lesson: a stage that fails silently is more dangerous than one that crashes loudly.*
2. **70 good clips were wrongly thrown out as “two speakers.”** The gender detector was labelling chunks of one expressive male storyteller as “female” (his pitch swings fooled it), which tripped the two-speaker check. I added a per-source `so1o: true` flag meaning “I listened, it’s one person – skip that check.”
3. **The LLM returned blank answers and burned through credits.** `sarvam-30b` is a *reasoning* model: left alone it spent its entire token budget “thinking” and returned nothing. Every call was both expensive and useless. The fix was one setting (`reasoning_effort=None`) that turns thinking off – about 60x cheaper and 30x faster.
4. **A whole source was silently skipped** because its YouTube id had an underscore in it, which broke my filename parsing. Fixed by matching against the known source names.
5. **Clips started too abruptly.** I was cutting flush to the first sound, which clipped the start of words. Added a small 120 ms cushion taken from the surrounding silence. For ASMR I also had to stop long silent gaps from being packed into clips.
6. **A source I labelled “English” was actually Telugu.** `En_Podcast_Raj` is Telugu with heavy English mixing, but I’d set it to English – so Sarvam faithfully wrote the Telugu in Latin letters

("O pakka divorce teesukunna vaadu..."). It *sounded* Telugu but *read* English. Only listening caught it. Relabelling it Telugu and re-running gave proper Telugu script, and I noticed the main speaker is female – which helped balance the dataset.

7. **Whisper can't be detected from text.** The LLM only sees the words, so it can't know a clip is whispered. For ASMR sources I just tell it directly with a `whisper: true` flag.
8. **Over-collect, then balance.** I gathered 617 clips, then wrote `balance.py` to trim down to an even spread across every source and both languages, always keeping any clip a human had already checked.

3. What worked and what didn't

Worked

- Cutting only at silences: all 190 clips land between 18 and 28 seconds, none over the limit, no chopped words.
- Sarvam's ASR is strong on clean audio – on the clips a human checked, the transcripts were good enough to accept as-is. It also handles Telugu-English mixing and writes the English in Telugu script.
- Diarization cleanly pulled single speakers out of 3-4-person podcasts.
- The crowd-noise filter removed 79 noisy stand-up clips.
- The diversity push worked: I went from an almost entirely "narrative + neutral" set to one with real conversational and playful clips, and from 90% male to ~40% female.

Didn't

- **Some crowd noise still slips through.** My noise check is a blunt "is more than 25% of the clip noisy?" rule, and the detector can't tell "clean voice" from "voice with light laughter underneath", so a few stand-up clips kept faint audience noise. This is the dataset's weakest spot.
- **English ASMR produced zero clips** – it's whispered so quietly that my silence detector treated the whole thing as silence.
- The audio-emotion model and the text-LLM only agree ~15% of the time on Telugu, because that model was trained on English/German. I use it only to *flag* clips for human attention, never as a label.
- The set still leans narrative – long-form Indian content is mostly monologue.

4. Quality observations and decisions I made

The final set, by the numbers:

Language	Telugu 39.7 min (95 clips) · English 39.2 min (95)
Gender	male 41 min · female 31 min (~40%) · unknown 7
Whisper clips	5

Collected -> kept

617 -> 190 (dropped: 60 music, 79 crowd, 13 by hand, 275 to balance)




Figure 2: Minutes per source. The over-collect-then-balance step spreads the time evenly across all 14 sources, so no single channel dominates.




Figure 3: Emotion spread. Adding stand-up, podcasts and ASMR widened this from a neutral-only spike to nine emotions.

On accuracy (WER/CER). Measured on the 37 human-checked clips, error rate is 0.00 – but I want to be honest about what that means: the reviewers *accepted Sarvam’s transcripts unchanged*, so this is really an “approval rate” on clean clips, not a strict error rate against a from-scratch transcription. It’s a genuinely good sign for Sarvam on clean audio, but a stricter test (re-transcribing blind) would give a truer number.




Figure 4: Speaking style. Still narrative-heavy, but conversational is now well represented.

Decisions worth calling out:

- **Two opinions on emotion.** The text LLM reads the words; a separate model listens to the sound. When they disagree (they did on 161 of 190 clips), that clip goes to the front of the human-review queue – that’s where a human’s time is worth the most.
- **Listening is the real work, so I made it easy.** I built a review app a non-technical person could use, and my Telugu-reading grandmother checked the Telugu clips – better ground truth than me, since I don’t read Telugu. She listened, confirmed or fixed the text, tapped an emoji for the feeling, and saved (or rejected bad clips). 37 verified, 13 rejected.




Figure 5: The review app: play the clip, fix the transcript if needed, tap how it sounds, Save or Reject. It opens on a small, balanced “review set” so you check a handful of important clips, not all 190.



Figure 6: The bilingual version made for my grandmother – Telugu labels, big text, emoji for emotion so nothing needs reading in English, and a done/left counter.

5. What I'd improve given more time

First, how I worked. My biggest mistake was process, not code: once the first small test worked, I built the whole pipeline in one go and leaned heavily on AI agents to write it – so the silent bugs above only surfaced when I ran everything together. Next time I'd test each stage on two or three clips and check its output myself before moving to the next, and treat the AI as a pair-programmer I review rather than letting it drive.

Then, the data:

- **Better noise detection** (the top fix). Instead of one blunt threshold, measure each clip's signal-to-noise ratio, add a model that specifically detects applause/laughter/crowd, gently denoise only the borderline clips, and send more of them to a human ear.
- **Shorter clips** (say 8-12 seconds). Smaller windows span fewer speaker turns, so there's far less chance of two voices landing in one clip – the main risk with podcasts.
- **Per-source silence settings** so quiet ASMR segments properly (this would recover the English ASMR I lost).
- **A multilingual emotion model** that actually knows Indian languages, so the audio and text emotion opinions agree more often.
- **More female Telugu sources, more genres, and a YouTube scraper** to scale past a hand-picked roster – plus a fuller human-review pass for a larger trusted set.

Ethics and licensing

Every clip records where it came from and what kind of source it is. I avoided copyrighted film and music, and dropped anything with a music bed. I haven't individually cleared redistribution rights for the underlying audio, so the dataset is meant for **research and educational use**; the licence

tag covers the annotations I added. Anyone using it commercially should clear the source audio themselves.